

(1) Dataset loader

cases.jsonl (20 synthetic + 61 OncQA)



(2) Perturbation engine

Name-Bank x Geo substitution, SHA-256 keyed



(3) Model harness

OpenAI + Groq via stdlib urllib



(4) LLM-as-annotator

Llama-3.1-8B, JSON-constrained



(5) Metrics & stats

TSR, RCR, RCER, GDI; Wilcoxon; BCa bootstrap

Cross-cutting:

- Token-bucket rate limiter (per-model)
- SHA-256 idempotency cache
- Manifest hashing (dataset, Name-Bank)
- YAML-driven reproducibility